

6 Related Work

Several past studies have curated character attribution datasets for narrative understanding. Finlayson (2015) annotated seven character archetypes in Russian folktales to create the ProppLearner corpus. Skowron et al. (2016) labeled hero, antagonist, sidekick, spouse and supporting roles in action-genre movies. Brahman et al. (2021) collected character descriptions from online study guides such as LitChart and Sparknotes, and created the LiSCU dataset for the character identification and description generation task. Sang et al. (2022) curated MBTI personality types from the personality database for movie characters. Yu et al. (2023) annotated reader’s notes for character traits in Chinese-translated Gutenberg books. Baruah and Narayanan (2024) annotated screenplay excerpts for character attributes and evaluated in-context and chain-of-thought learning methods on the attribution task. Table 1 compares these datasets against CHATTER and CHATTEREVAL.

7 Conclusion

We proposed the CHATTER and CHATTEREVAL datasets for the narrative character attribution task. We addressed the limitations of previous datasets by curating a resource that is scalable, generalizable, well-defined and discrete. Experiments showed that CHATTER can serve as a reliable training set and CHATTEREVAL can be used as the evaluation benchmark for character attribution modeling. Future work includes developing character attribution models on our datasets to aid creators and writers in analyzing their narratives.

8 Limitations

We define the character attribution task as a binary classification task, where, given the character-trope pair and the screenplays of the movies in which the character appeared, the model should predict whether or not the character portrayed the trope. This formulation has some limitations. First, the screenplay’s narrative may not exactly resemble the story told by the movie because of tweaks made during the filming process. Publicly available screenplays are rarely the final script but drafts from earlier in the production stage. Second, movies could have visual cues like the non-verbal behavior of the character that are missed by the screenplay text. Lastly, a character can appear in multiple movies, and our dataset does not

contain scripts for all of them. These limitations have important implications because the TVTropes contributors and our raters draw their knowledge of the character from all sources. In contrast, the attribution model predicts the attribute labels based only on the screenplay. There could be character-attribute pairs for which the model has insufficient information and could lead to poor recall, as shown by the performance of the models in Table 5.

References

- Marah Abdin, Jyoti Aneja, Hany Awadalla, Ahmed Awadallah, Ammar Ahmad Awan, Nguyen Bach, Amit Bahree, Arash Bakhtiari, Jianmin Bao, Harkirat Behl, et al. 2024. Phi-3 technical report: A highly capable language model locally on your phone. *arXiv preprint arXiv:2404.14219*.
- Mahmoud Azab, Noriyuki Kojima, Jia Deng, and Rada Mihalcea. 2019. [Representing movie characters in dialogues](#). In *Proceedings of the 23rd Conference on Computational Natural Language Learning (CoNLL)*, pages 99–109, Hong Kong, China. Association for Computational Linguistics.
- David Bamman, Brendan O’Connor, and Noah A. Smith. 2013. [Learning latent personas of film characters](#). In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 352–361, Sofia, Bulgaria. Association for Computational Linguistics.
- David Bamman, Ted Underwood, and Noah A. Smith. 2014. [A Bayesian mixed effects model of literary character](#). In *Proceedings of the 52nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 370–379, Baltimore, Maryland. Association for Computational Linguistics.
- Sabyasachee Baruah and Shrikanth Narayanan. 2023. [Character coreference resolution in movie screenplays](#). In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 10300–10313, Toronto, Canada. Association for Computational Linguistics.
- Sabyasachee Baruah and Shrikanth Narayanan. 2024. [Character attribute extraction from movie scripts using llms](#). In *ICASSP 2024 - 2024 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 8270–8275.
- Faeze Brahman, Meng Huang, Oyvind Tafjord, Chao Zhao, Mrinmaya Sachan, and Snigdha Chaturvedi. 2021. [“let your characters tell their story”: A dataset for character-centric narrative understanding](#). In *Findings of the Association for Computational Linguistics: EMNLP 2021*, pages 1734–1752, Punta Cana, Dominican Republic. Association for Computational Linguistics.

- Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.
- Mark A. Finlayson. 2015. [ProppLearner: Deeply annotating a corpus of Russian folktales to enable the machine learning of a Russian formalist theory](#). *Digital Scholarship in the Humanities*, 32(2):284–300.
- Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. 2024. Gpt-4o system card. *arXiv preprint arXiv:2410.21276*.
- Naoya Inoue, Charuta Pethe, Allen Kim, and Steven Skiena. 2022. [Learning and evaluating character representations in novels](#). In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 1008–1019, Dublin, Ireland. Association for Computational Linguistics.
- Albert Q Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, et al. 2023. Mistral 7b. *arXiv preprint arXiv:2310.06825*.
- Carl Gustav Jung. 2014. *The archetypes and the collective unconscious*. Routledge.
- Vincent Labatut and Xavier Bost. 2019. Extraction and analysis of fictional character networks: A survey. *ACM Computing Surveys (CSUR)*, 52(5):1–40.
- Dawei Li, Hengyuan Zhang, Yanran Li, and Shiping Yang. 2023. [Multi-level contrastive learning for script-based character understanding](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 5995–6013, Singapore. Association for Computational Linguistics.
- Andrew Piper, Richard Jean So, and David Bamman. 2021. [Narrative theory for computational narrative understanding](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 298–311, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Vladimir Propp. 1968. Morphology of the folktale. *U of Texas P*.
- Machel Reid, Nikolay Savinov, Denis Teplyashin, Dmitry Lepikhin, Timothy Lillicrap, Jean-baptiste Alayrac, Radu Soricut, Angeliki Lazaridou, Orhan Firat, Julian Schrittwieser, et al. 2024. Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context. *arXiv preprint arXiv:2403.05530*.
- Yisi Sang, Xiangyang Mou, Mo Yu, Dakuo Wang, Jing Li, and Jeffrey Stanton. 2022. [MBTI personality prediction for fictional characters using movie scripts](#). In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pages 6715–6724, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.
- Marcin Skowron, Martin Trapp, Sabine Payr, and Robert Trappl. 2016. Automatic identification of character types from film dialogs. *Applied Artificial Intelligence*, 30(10):942–973.
- Zhilin Wang, Weizhe Lin, and Xiaodong Wu. 2021. [Learning similarity between movie characters and its potential implications on understanding human experiences](#). In *Proceedings of the Third Workshop on Narrative Understanding*, pages 24–35, Virtual. Association for Computational Linguistics.
- Mo Yu, Jiangnan Li, Shunyu Yao, Wenjie Pang, Xiaochen Zhou, Zhou Xiao, Fandong Meng, and Jie Zhou. 2023. [Personality understanding of fictional characters during book reading](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 14784–14802, Toronto, Canada. Association for Computational Linguistics.
- Lisa M Pytlik Zillig, Scott H Hemenover, and Richard A Dienstbier. 2002. What do we assess when we assess a big 5 trait? a content analysis of the affective, behavioral, and cognitive processes represented in big 5 personality inventories. *Personality and Social Psychology Bulletin*, 28(6):847–858.

A Annotation

We give the following annotation instructions to the Amazon MTurk workers.

A trope is a storytelling device or convention the storyteller uses to describe situations the audience can easily recognize, often used to stereotype characters. Think of them as character attributes.

Your task is as follows:

1. Identify the movie character from the pictures and the given movie(s). If you are having trouble recognizing the character, click on the Character Page link(s).
2. Read the definition of the character trope. You can open the link to know more.
3. Choose yes, maybe yes, not sure, maybe no, or no to answer whether the character portrays or is associated with the trope.
4. (Optional) Explain your choice in the provided text area or leave any other comment behind.

Figure 1 shows the annotation interface that the workers use to label character attribution.

B TVTropes

TVTropes is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 3.0 Unported License. This license allows for the copy and redistribution of the material in any medium or format for non-commercial purposes. CHATTER is published for academic research purposes and does not infringe the TVTropes license.

C ScriptsonScreen

ScriptsonScreen aggregates movie scripts from different websites such as IMSDB⁸, Dailyscripts⁹ and Script Slug¹⁰. All the movie scripts we use are

⁸<https://imsdb.com/>

⁹<https://www.dailyscript.com/>

¹⁰<https://www.scriptsbug.com/>

licensed for fair use and available for public download.

D TROPES

There are 445 tropes in total that have at least one character portraying them in our dataset. We could not figure out a systematic way to cluster them because the tropes are very different from each other. Therefore, we simply list the most common tropes and their abbreviated definition. Tables 6 and 7 list the tropes that are portrayed by at least two characters in our dataset.

E Prompt

We used the following prompt template in our experiments.

A character trope is a story-telling device used by the writer to describe characters.

Given below is the definition of the \$TROPE\$ trope enclosed between the tags <TROPE> and </TROPE>.

Following that is a movie script enclosed between the tags <SCRIPT> and </SCRIPT>. The character "\$CHARACTER\$" appears in the movie script.

Read the movie script carefully and based on that answer yes or no if the character "\$CHARACTER\$" portrays or is associated with the \$TROPE\$ trope. If yes, give a brief explanation. Answer based only on the movie script. Do not rely on your prior knowledge.

<TROPE>
\$DEFINITION\$
</TROPE>

<SCRIPT>
\$SCRIPT\$
</SCRIPT>

Does the character "\$CHARACTER\$" portray or is associated with the \$TROPE\$ trope in the above movie script?
Answer yes or no.